



Gautam Ethiraj

Follow

Feb 12, 2022 · 3 min read · Listen



Tien Nguyen

nguyen.ttq.tien@gmail.com

Continue as Tien



Save



What is nn.Embedding really?

$$\begin{array}{c}
 \begin{bmatrix} 0 & 0 & 0 & 1 & 0 \end{bmatrix} \\
 \text{One-hot vector}
 \end{array}
 \times
 \begin{array}{c}
 \begin{bmatrix} 8 & 2 & 1 & 9 \\ 6 & 5 & 4 & 0 \\ 7 & 1 & 6 & 2 \\ 1 & 3 & 5 & 8 \\ 0 & 4 & 9 & 1 \end{bmatrix} \\
 \text{Embedding Weight Matrix}
 \end{array}
 =
 \begin{array}{c}
 \begin{bmatrix} 1 & 3 & 5 & 8 \end{bmatrix} \\
 \text{Hidden layer output}
 \end{array}$$

In this brief article I will show how an embedding layer is equivalent to a linear layer (without the bias term) through a simple example in PyTorch. This might be helpful getting to grips with the nitty-grittys of implementing it in your models (even if you might have already *conceptually* known about this equivalence).

Embedding vs Linear definition-wise?

An embedding is basically the same thing as a linear layer but works differently in that it does a lookup instead of a matrix-vector multiplication.

Why use an embedding when we have a linear layer?

An embedding is an efficient alternative to a single linear layer when one has a *large number of input features*. This may happen in natural language processing (NLP) when

one is working with text data or in some (language-like) tabular data that is treated as a bag-of-words (BoW). In such cases its also quite common to have the input data available as a sparse matrix (typically a result of an output from sklearn's CountVectorizer or TfidfVectorizer as a *sparse.scipy.csr_matrix*) and it is memory-inefficient to convert that in to a dense matrix but really easy to access its non-zero elements and their positions directly instead (using the *data* and *indices* attributes).

Lets look at the example

Assume *X_train* is the input data of the training set in a sparse matrix format. For this simple example I am creating a small sparse matrix so that one can easily follow along. I also extract one *row* of this matrix which is like taking one sample from the dataset that would go through the forward method of your model. The *.getrow()* method helps here and that's what you will need in your PyTorch *Dataset/Dataloaders*.

```
# Initialize a sparse matrix: This could be your training set
X_train = csr_matrix(np.array([[1, 0, 1, 0],
                               [0, 0, 1, 1],
                               [1, 1, 1, 0]]))

# Get one row: One sample in the training set
row = X_train.getrow(0)
```

Now let's pass the training example *row* through the linear layer and the embedding so that we get the same result in each case.

nn.Linear

```

w_linear = nn.Linear(4, 3, bias=False)
w_linear.weight

Parameter containing:
tensor([[ 0.3444,  0.2103, -0.2834,  0.4480],
        [-0.2406,  0.2936,  0.1943,  0.0454],
        [ 0.3183, -0.2030, -0.1320, -0.1592]], requires_grad=True)

print(w_linear(torch.FloatTensor(row.toarray()))) row is now a dense matrix
tensor([[ 0.0609, -0.0463,  0.1862]], grad_fn=<MmBackward>)

```

nn.Embedding

```

w_embedding = nn.Embedding(4, 3).from_pretrained(w_linear.weight.T)
w_embedding.weight

```

Transposed!

Open in app ↗

Sign up

Sign In



Search Medium



nn.Embedding

```

print(w_embedding(torch.tensor(row.indices)).sum(0))
tensor([ 0.0609, -0.0463,  0.1862])

```

Same final result with an embedding layer as with a linear layer!

The outputs are the same. Yay! A couple of observations to keep in mind when you're using this in your own *nn.Module*:

1. The embedding weights and the linear layers weights are transposed to each other.
2. The linear layer *w_linear* does the actual matrix vector multiplication and therefore needs the row to be converted to dense format. In contrast, *w_embedding* just needed the indices of *row* to do a lookup. Not only is this faster, but it's also quite convenient with the *scipy.sparse.indices* attribute that is available for the sparse matrix!
3. The embedding requires the *sum(0)*. Don't forget it!

Thanks Jeremy Howard for exposing me to this idea several years ago!

Pytorch

Embedding

Linear

Logistic Regression

Optimization

[About](#) [Help](#) [Terms](#) [Privacy](#)

Get the Medium app

